

Allianz Data Science Challenge



A short description of the project.

Project Organization

```

├── LICENSE                <- Open-source license if one is chosen
├── Makefile               <- Makefile with convenience commands like `make data` or
`make train`
├── README.md             <- The top-level README for developers using this project.
├── data
│   ├── external          <- Data from third party sources.
│   ├── interim           <- Intermediate data that has been transformed.
│   ├── processed         <- The final, canonical data sets for modeling.
│   └── raw               <- The original, immutable data dump.
├── docs                  <- A default mkdocs project; see www.mkdocs.org for details
├── models                <- Trained and serialized models, model predictions, or
model summaries
├── notebooks             <- Jupyter notebooks. Naming convention is a number (for
ordering),
                        <- the creator's initials, and a short `-` delimited
description, e.g.
                        <- `1.0-jqp-initial-data-exploration`.
├── pyproject.toml        <- Project configuration file with package metadata for
tools like black
├── references            <- Data dictionaries, manuals, and all other explanatory
materials.
├── reports
│   └── figures           <- Generated graphics and figures to be used in reporting
├── requirements.txt      <- The requirements file for reproducing the analysis
environment, e.g.
                        <- generated with `pip freeze > requirements.txt`
├── setup.cfg            <- Configuration file for flake8
└── allianz_data_science_challenge <- Source code for use in this project.
    ├── __init__.py      <- Makes allianz_data_science_challenge a Python
module
    |

```

```
├── config.py          <- Store useful variables and configuration
├── dataset.py         <- Scripts to download or generate data
├── features.py        <- Code to create features for modeling
├── modeling
│   ├── __init__.py
│   ├── predict.py     <- Code to run model inference with trained models
│   └── train.py       <- Code to train models
```

How the Project Works (App Execution Flow)

The heart of this project lies in the `app.py` script, which serves as the main entry point for training the model and making predictions. It leverages modular components from the `allianz_data_science_challenge` package.

File: `app.py`

This file acts as a **command-line interface (CLI)** for:

1. Loading a sample customer profile.
 2. Making a prediction using the current best model.
 3. Automatically training the model if no valid model exists.
 4. Printing a human-readable result.
-

Workflow Summary

1. Input Preparation:

- A sample customer record is manually defined (you can replace it with dynamic input later).
- The record includes key features like `age`, `job`, `contact`, `month`, `duration`, and economic indicators.

2. Prediction Attempt:

- The sample data is passed to the `predict()` function.
- If the model is already trained and valid, it returns a binary prediction (0 = No, 1 = Yes).

3. Model Retraining if Needed:

- If prediction fails (e.g., missing or corrupted model file), the script falls back to `train_model()` from the `train.py` module.
- After successful training, it retries the prediction with the freshly trained model.

4. Output:

- Displays whether the customer is predicted to subscribe to a term deposit.

- Handles errors gracefully and provides user-friendly feedback.

🗄️ Key Modules Used

Module	Purpose
<code>train_model()</code>	Trains and saves the best model.
<code>predict(df)</code>	Loads the model and predicts on new data.
<code>config.MODELS_DIR</code>	Manages file paths like <code>best_model.pkl</code> .

📦 Future Improvements

- Replace hardcoded input with CLI flags, web input, or batch prediction.
 - Add API or Web UI on top (e.g., Flask/FastAPI).
 - Include logging, monitoring, and metrics tracking for production usage.
-

🔧 Example CLI Output

```
Bank Marketing Model - Prediction
=====
Input data: {...}
Prediction: 1
Prediction is 1.
This client is predicted to subscribe to a term deposit.
```

🎯 Problem Framing

The core business challenge is **optimizing the cost-efficiency** of a direct call campaign to sell term deposit products to existing bank customers.

- **Objective:** Reduce the number of unnecessary calls made to customers who are unlikely to subscribe, while retaining as much business value (conversions) as possible.
- **Constraints:**
 - A single call costs approximately **8€**.
 - Each successful subscription yields **~80€** in profit.
 - The **only channel** for conversion is the phone call — no call, no sale.
 - If a customer does **not** subscribe during the call, they are considered permanently uninterested.
- **Framed as a supervised classification problem:**
 - Input: Customer features (demographics, contact history, previous campaign outcome, etc.)
 - Output: Binary label — **1** if the customer subscribed, **0** otherwise.
- **Business metric focus:**
 - Maximize profit = $(80€ \times \text{true positives}) - (8€ \times \text{total calls})$

- Optimize for **balanced F1 score** and **profit gain** rather than raw accuracy, due to class imbalance (conversion rate is low).

Key Drivers of Conversion (Based on Model Insights)

1. Month of Contact

- Conversion rates were **significantly higher** in **March, December, and October**.
- Likely due to seasonal behaviors or financial planning cycles.

2. Outcome of Previous Campaign (**poutcome**)

- Customers with a **previous successful interaction** had a much **higher likelihood of converting** again.
- Indicates strong behavioral momentum or positive prior engagement.

3. Type of Contact

- **Cellular** contact showed better conversion than **telephone** (landline).
- Suggests that mobile outreach may reach customers at more convenient times or signals better accessibility.

4. Job Type

- **Students and retirees** had notably **higher conversion rates**.
- Indicates they might be more receptive to long-term financial planning products.

5. Education Level

- Customers with **tertiary education** were more likely to subscribe.
- Suggests education level correlates with product understanding or trust.

These findings help build a data-driven approach to **select weekly cohorts of customers** who are:

- Most likely to convert based on these predictive features.
- Actionable for real-time prioritization and cost-saving contact decisions.

Business Metrics

- **Profit if calling ALL customers:** €8,896.00
- **Profit using model predictions:** €21,016.00
- **Improvement from model:** €12,120.00
- **Cost savings from avoided calls:** €56,040.00
- **Total business value delivered:** €77,056.00
- **Maximum possible profit:** €74,800.00
- **Profit efficiency:** 28.1%